Working with Claude

Set up an AI assistant that remembers your work between sessions, files things for you, and checks its own output. For a book, a business, a website, a family history: anything.

The idea

Most people use AI as a chat window: you ask, it answers, and next week it has forgotten everything. This setup is different. You give Claude **its own folder on your computer**, like a filing cabinet, with one folder per project. In it, Claude keeps notes on your project, how you like things done, what you've decided and why, and a plan for next time. Each time you open the folder, Claude reads its notes first and picks up where you left off.

> **Each time you come back, once you're set up**
> 1. Open the Claude app → **Code** → **+ New session** → choose your project folder.
> 2. Say **hello**. Claude offers **one small next step**, already worked out. Say **"go"**.
> 3. Work by **talking**. Put new files in the **inbox** folder and say "process the inbox".
> 4. Stop whenever you like. Say **"that's it for today"** and Claude tidies up and saves, asking you nothing (for assessed work with a logbook, one question). Or just close the window: nothing's lost.

> **Already comfortable with a terminal?** Skip to Part 3 for the terminal install (it also covers Linux), then come back for Step 3 (open `claude` in your **Claude projects** folder and paste the sentence). Setup is the same: `/setup`. If your work is marked or assessed (a dissertation, coursework), read "Assessed work" in Part 2 before you start.

PART 1 · GETTING STARTED

Installing takes a little while. After that, setup takes about ten minutes and ends with a real piece of work done.

What you need

- A Mac or Windows computer.
- A **paid Claude plan** (Pro or above), a monthly subscription you can cancel. The free plan doesn't include what you need. Prices are at *claude.com/pricing*.
- Nothing else to download: in Step 3, Claude fetches the **starter kit** for you from *github.com/teedor/ai-workspace-kit*.
- **Optional, for later:** two or three examples of your own work to show Claude your style, like an email you were pleased with, an old menu, a chapter or a letter.

Paid plans come with a usage allowance. If Claude says you've **reached your limit**, nothing is broken. It tells you when the allowance resets, and you won't be charged extra unless you've switched on "extra usage" in your account.

STEP 1

Install the Claude app

1. Go to **claude.com/download**, download the app for your computer, and install it like any other app.
2. Open it and sign in.
3. **Privacy:** in the app's Settings, under Privacy, decide whether "Help improve Claude" should be on. Turn it off if you don't want your work used to train future models.

STEP 2

Git (Windows only; Mac users skip this)

Git is a free tool that saves snapshots of your work, so you can go back. Claude runs it for you.

- **Windows:** go to **git-scm.com**, download Git for Windows and install it. You'll see a dozen screens full of technical words. Ignore them all and keep clicking **Next**; the defaults are right. When it's finished, **close the Claude app completely and open it again**, so it notices Git.
- **Mac:** nothing to do now. The first time Claude uses Git (in Step 3), a box may pop up offering to install "command line developer tools". Click **Install**, wait until it finishes (it can take a while), then tell Claude "done". If nothing seems to happen, look behind the Claude window.

STEP 3

Get the kit and make your first project folder

1. **Open your home folder.** Mac: in Finder, choose **Go → Home**. Windows: in File Explorer, click the address bar, type `%USERPROFILE%`, and press Enter. (Avoid Documents and Desktop. They're often synced to iCloud or OneDrive, and syncing can scramble the snapshots.)
2. Make a new folder there called **Claude projects**. All your projects will live in it. Leave it empty.
3. In the Claude app, click the **Code** tab at the top. Despite the name, it's for any kind of work.
4. Click **+ New session** in the left-hand list, or press `Cmd + N` (Mac) or `Ctrl + N` (Windows).
5. Around the box where you type, at the bottom, are a few settings. Set:
   - **Environment: Local** (it means "on this computer").
   - **Project folder** (it may say **Select folder**): click it and choose **Claude projects**. To reach your home folder in that window: on a Mac, press `Cmd + Shift + H`; on Windows, click the address bar at the top, type `%USERPROFILE%` and press Enter. If the app asks whether you trust this folder, say yes: it's your own.
   - **Permission mode: Manual** (older versions say *Ask permissions*). Claude then asks before every change.

   The app gets updated, so it may look slightly different, but those choices will be there.
6. Copy this sentence, paste it into the typing box, and press **Enter**. Paste it rather than dictating it: the web address has to be exact. (Line breaks don't matter.)

```
Set me up with the AI Workspace Kit: clone github.com/teedor/ai-workspace-kit
into ~/.claude/workspace-kit (or update it if it's already there, deleting
nothing), then follow its INSTALL.md; if Git is missing, stop and tell me how
to install it.
```

7. Claude asks what your project is called, makes a **project folder** for it inside **Claude projects** (for example *Bakery* or *My novel*), and tells you how to open it.

Expect several permission requests, and Claude says what each one is for. The first saves the kit in Claude's own settings folder (`.claude` in your home folder); Claude always asks before writing there. Choose the plain **Yes** each time, not an option that stops it asking. If Claude gives you commands to type yourself, you're in the **Chat** tab: switch to **Code** and start again.

Only ever paste a "fetch this and follow it" sentence from somewhere you trust, like this guide. Claude follows what it finds there, within the permissions you give it.

> **Can't download the kit?** Some work laptops block it, and if a GitHub sign-in box pops up, close it: the kit doesn't need an account. Instead, go to *github.com/teedor/ai-workspace-kit*, click the green **Code** button → **Download ZIP**, and unzip it where it lands, in Downloads (Mac: double-click it; Windows: **right-click** → **Extract All…** → **Extract**). Then paste this instead: *"Set me up with the AI Workspace Kit from the ZIP I've unzipped in my Downloads folder: move the kit folder (the one with INSTALL.md in it) to ~/.claude/workspace-kit, deleting nothing, then follow its INSTALL.md."*

STEP 4

Your first session

Claude works on **one project folder at a time**, and you choose it at the start of each session. No rush, and you can stop halfway: setup picks up where it left off.

1. Start a **+ New session** as in Step 3, with the same settings, but this time choose your **project folder** (*Bakery*), not *Claude projects*. (It has to be a new session: Claude only reads a project's notes when the session is opened on that folder.)
2. Type `hello` and press **Enter**.
3. **Check the reply.** It should start by naming your folder, like *"📁 You're in Bakery."* If it doesn't, ask *"Which folder are you working in?"* If the answer isn't your project folder, start a new session and choose again.
4. Type `/setup` and press **Enter**. (Typing `/` pops up a menu; ignore it and keep typing.)

> **How to tell you've got the right folder:** your project folder has a file called `CLAUDE.md` and a folder called `inbox` directly inside it.

During setup

Claude asks a handful of questions, one at a time. Answer by talking or typing. It sets up the snapshots, then asks:

- to confirm the project's **name** (it suggests the folder name), what it is, and whether it involves code;
- whether it's **marked or assessed work** (coursework, a dissertation) or has rules about using AI. If it is, Claude adds the study pack (see Part 2);
- whether you'll put in **real people's details or data**: emails, orders, invoices, family records, or datasets about people. Those all count, so say yes if you'll use any of them;

- for **one small thing** you could do on it right now. Then it helps you do it.

Everything else (a file about you that all your projects share, a style guide from your examples, dictation, Word) is offered **later, one thing at a time**, at the end of a session that went well.

Expect permission requests, some in technical words. A few are for things outside your project folder, such as saving your name for the snapshots. **Claude tells you before each of those that it's coming and why.** Allow them.

> **You're set up when** Claude says *"You're set up."* If something goes wrong partway through, say *"something went wrong, explain in plain English"*. Or start a new session in the same folder and type `/setup` again: it picks up where it left off.

Saying yes (or no) to Claude

Claude often pauses to ask permission, sometimes in words like *Bash* or *git commit*. That's normal.

- **Allow it** if it matches what you just asked for, or what Claude just told you it was about to do.
- **Say no**, or ask *"what does that do, in plain English?"*, if it's a surprise, especially anything about deleting, sending, publishing or money.

Coming back the next day

1. Claude app → **Code** → **+ New session**.
2. Choose your project folder (it'll be in the recent list), Local, Manual.
3. Say **hello** (or `/start`). Check the 📁 line names the right folder.
4. Claude offers one next step. Say **"go"**, **"something else"**, **"I've got longer"** or **"low day"**.

Closed the window without finishing? Been away for weeks? Nothing is lost, and Claude doesn't make a thing of it. It catches up quietly and offers a small way back in.

Short sessions, long sessions, low days

Sessions can be as short or as long as you like. Claude doesn't ask you to commit to a length. It learns how you tend to work from how your sessions actually go, and sizes the next step to fit. When you're on a roll, it offers the next step. When you've done one and want to stop, stop.

On a day when you've got nothing in you, say **"low day"**. Claude offers three tiny things: talk for two minutes about what's on your mind, just open the file and read it, or one easy step. Or close it: the next step is written down for whenever you're back. There are no streaks and no "days since". It only keeps a record of what got done.

Talk, don't type

The biggest tip in this guide. Use your computer's dictation for almost everything you say to Claude. Most people speak about three times faster than they type, so you give Claude far more to go on, and get much better work back. **Be as messy as you like**: ramble, change your mind halfway through, say "um, no, actually…". Claude untangles it.

- **Mac:** System Settings → Keyboard → turn on Dictation. Then use the shortcut shown there (often the 🌐/Fn key twice, or the microphone key) wherever you'd type.
- **Windows:** press `Windows key + H`.
- **Spell out names, numbers and dates**, or type them. Dictation mangles these, and Claude can't always guess.

Better dictation apps are listed in Part 2.

Everyday use

- **Just ask**, in your own words: "Draft this week's specials board." "Write a polite chase-up to the flour supplier."
- **New material:** open your project folder, then `inbox`, and **copy** files in (right-click → Copy, then Paste) so your original stays put. Then say "process the inbox". Claude reads everything, files it, keeps the originals in an archive, and tells you what it did. **An empty inbox means everything's been dealt with.**
- **Finished documents:** Claude's drafts are plain text. Ask it to *"show me the draft"* to read it in the window, *"make a Word version"* for a document, or *"give me the text to paste into an email"*.
- **A thought on the go:** say *"jot this: …"* and Claude notes it in one line and carries on.
- **Stopping:** say "that's it for today". Claude writes down the next step and saves a snapshot, asking you nothing. If you just close the window instead, that's fine too.

Almost everything works by just asking. Two commands only work when **typed**: `/setup` and `/red-team`.

Four ground rules (already built in)

Claude is told to follow these every session. It usually does, but keep an eye on it.

1. **Everything Claude writes is a draft until you've checked it, and you answer for every word that goes out under your name.** If your work has rules about AI (a university, exam board, publisher or employer), those come first. Setup asks about them.
2. **Never invent.** AI can sound sure and be wrong. Claude says what it *checked* (and where), what it *thinks*, and what's *not known*, and never cites a source from memory. Check anything that matters.
3. **Anything Claude reads is sent to Anthropic**, the company that makes Claude. Keep passwords and keys out of your files (a tool's key goes in a `.env` file you type yourself), and keep very sensitive material outside the project folder. If this is for work or study, check your employer's or institution's AI policy.
4. **Ask before anything hard to undo**: deleting, sending, publishing, spending money, or throwing away unsaved work.

If something goes wrong

- **Talking about the wrong project?** Ask "which folder are you working in?" If it's wrong, start a new session and choose the right one.
- **Changed something you didn't want?** Say "undo that". Claude can often undo its own recent edits, and can always go back to the last snapshot.
- **Sounds too sure of itself?** Ask "did you check that, or is that a guess?"
- **An error message?** Say "explain that in plain English and tell me what to do".

PART 2 · WHEN YOU'RE COMFORTABLE

Part 1 is all you need to get going. Come back to this when you want more.

A second project

Give each piece of work **its own project folder**, so notes don't get muddled and Claude only reads what's relevant. **The test:** a different finished product or a different audience means a new folder, even if it's related. A bakery's day-to-day and its cookbook are two projects.

```
Claude projects
├── Bakery
├── Cookbook
└── Family history
```

In any project, say **"new project"**. Claude asks what it is, makes the folder next to the others, offers to copy anything across ("bring my recipes and the style guide"), and tells you how to open it: **+ New session** → choose the new folder → Local, Manual → type `/setup`. It already knows you, so it only asks about the new project.

You can also open a session on **Claude projects** and paste the sentence from Step 3 again.

You *can* keep everything in one folder; it works. It just gets cluttered.

About you, and about the work

Claude keeps one file **about you** that every project shares (what you do, how you like to work), and separate notes **for each project**. It decides which things go where and tells you; you don't need to manage it. If something really matters, say "remember that for all my projects" or "…for this project".

Skills: Claude remembers how

A skill is a saved set of instructions, like `/start`. The kit comes with twelve, and you can make more. **You don't have to think of them yourself.** When you've asked for the same thing twice, or something works brilliantly, Claude offers, in one line:

```
Want me to remember how to do this, so next time you can just say
"chase the supplier"?
```

Say yes and it's done. Claude decides whether it's for this project or all of them. Say no and it won't ask again. You can also ask outright: "can you remember how to do that?"

Meetings

Got a meeting with a supervisor, manager, client or editor? Say so. Claude builds a one-page update from what you've done since last time, with your questions in your words. Afterwards, dictate your notes: Claude files them faithfully, puts actions and dates in your status, and records anything agreed.

Assessed work (students)

If setup hears that your work is marked or assessed, it adds a **study pack**. It turns the usual arrangement round: **you write, and Claude teaches and questions.**

- **Your university's rules come first.** Policies differ between universities and between modules. Your first step is to put your module's AI rules into `notes/ai-policy.md`, and to show your supervisor the "Rules for this work" section of `CLAUDE.md`, ideally at your next meeting.
- **Claude won't write assessed text or project code for you to submit.** It explains, asks questions, reviews what you've written, and helps you find bugs in your own code. Dictation for your report comes back as a faithful transcript, not tidied prose. If you insist, it's your call, and it's logged.
- **Everything is logged** in `notes/ai-use-log.md`, so your AI-use declaration is easy to write and accurate. Your git history, decisions log and experiment log show how the work developed.
- **Viva practice:** say "quiz me on this". Claude asks examiner-style questions, one at a time. Assessed work isn't signed off until you can explain it.
- **Honest references:** a sources register where only *you* mark a paper as read, and nothing unread gets cited.
- **For code and ML projects:** an experiment log with predictions written *before* each run, numbers only ever read from output files, a rule that the test set is touched once, and a red-team "evaluation auditor" that checks for data leakage and missing baselines.
- **If things get hard,** Claude will occasionally and briefly point you to a specific, practical route, like a two-line email to your supervisor or your university's support services. It won't nag, diagnose, or count the days.

Polishing work before it goes out

Say	What happens
"make it plainer"	**De-jargoning.** Claude lists hard words, long sentences and assumed knowledge, then rewrites in plain English, without making anything sound more certain than it is.
`/red-team` *(typed)*	**Red-teaming** means attacking your own work before anyone else does. Claude sends it to three separate critics who haven't seen your conversation (say, a regular customer, a fussy chef and a fact-checker), sifts their points, and gives you one list: must fix, should fix, optional. It uses a fair bit of your allowance, so save it for things that matter.
"let's sign things off"	Claude walks you through work waiting for your approval, one item at a time, showing exactly what needs checking.

Ideas and decisions

- **"I've had an idea…"**: Claude writes it up with the honest case *against* it and a test that would settle it, so a good-sounding idea doesn't slip into finished work untested.
- **"Let's go with…"**: Claude records the decision and why. If it replaces an older one, the old one is marked "replaced", never deleted, so you can see why you changed your mind.

Better dictation apps

App	What it is
Wispr Flow *wisprflow.ai*	Popular and polished. Hold a key, talk, and tidy text appears in any app. Mac, Windows, iPhone. Free up to 2,000 words a week; paid for unlimited. Works in the cloud, so check its privacy settings.
Superwhisper *superwhisper.com*	Mac, Windows, iPhone. Free and paid versions. Some of its speech models run entirely on your computer, so your voice never leaves it; check which one you're using.

Windows' built-in voice typing also uses an online service. Whichever you use, your words reach Claude as text anyway.

What's in a project folder

```
status.md     Where things stand: current goal, next steps, dates.
handover.md   The next step, ready for when you're back.
done.md       Everything finished. Nothing is ever removed.
inbox         Put things here unsorted. Claude sorts them and empties it.
notes         What Claude knows: style guide, glossary, open questions.
ideas         Untested ideas, with the case against.
decisions     What you decided, why, and what it replaced.
drafts        Work in progress.
finished      Work you've approved.
private       Confidential material and archived originals. Not in snapshots.
CLAUDE.md     Claude's standing orders for this project (read every session).
```

Some setup files are hidden on purpose. You never need to touch them.

Snapshots and backups

- **Snapshots** are saved every time you finish a session and every time you sign something off. Ask in plain English: "what changed last week?", "show me the menu as it was on Tuesday".
- **Not in snapshots:** the `private` folder, including the archive of your inbox originals. **Back up your whole Claude projects folder yourself**: Time Machine on a Mac, File History on Windows, or a copy on an external drive.

PART 3 · FOR THE TECHNICALLY MINDED

If you already use Claude Code

- When you accept the "About me" extra (offered after a session that went well), Claude adds a section to `~/.claude/CLAUDE.md`, which applies in **every** project, including existing repos. It shows you the text first and never replaces existing content.
- A personal skill in `~/.claude/skills/` overrides a project skill of the same name. Claude checks for clashes with the kit's skills when it sets up "About me", and offers to rename.
- The kit is kept as a git clone in `~/.claude/workspace-kit/`; its `claude-workspace` folder is the blank template. "New project" updates it (`git pull --ff-only`) and copies only that folder, so every project starts with its own fresh history. Updating the kit never changes projects you already have.

If you write code

- **Permission mode:** once you trust it, **Accept edits** mode cuts the prompts a lot. The kit already lets Claude save snapshots and update its own notes without asking, and still asks before `git push`, `rm` and anything that throws work away.
- **Big files and keys:** before every snapshot, Claude checks for files over 20 MB, folders of hundreds of new files (a dataset, say), and API keys written into code or notebooks. None of those get saved. A key that was ever committed or pasted into chat must be cancelled with the provider and replaced.
- **Existing project?** Don't copy the kit over it. Open your project and say: *"Bring in the workspace kit from ~/.claude/workspace-kit/claude-workspace (clone github.com/teedor/ai-workspace-kit there first if it's missing). Don't overwrite anything: merge CLAUDE.md, append to .gitignore, show me the changes first."* Then `/setup`.
- **One person per folder.** In a team repo, keep personal rules in `CLAUDE.local.md`, and agree with the team before committing workspace files.
- Code keeps its own layout. "Handover" commits only workspace files. "Sign off" runs your tests, commits reviewed code, and asks before committing to `main`. The red team's checker runs code only against test data.
- `.claude/settings.json` blocks Claude's file tools from reading `.env` and `.pem` files (including `.env.example`), and makes it ask before `git push`, `rm`, `git reset`, `git checkout` and similar, in Bash and PowerShell. It sets the starting permission mode to asking. These are guard rails, not a sandbox.

Prefer the terminal?

The app's Code tab and the terminal are the same Claude Code: same folders, skills and rules.

1. **Open a terminal.** Mac: `Cmd + Space`, type `Terminal`, press Enter. Windows: `Windows key + X`, then **Terminal** (Windows 11) or **Windows PowerShell** (Windows 10). Linux: your usual terminal; use the

Mac command below.

2. **Install (once):**

   ```
   Mac:      curl -fsSL https://claude.ai/install.sh | bash
   Windows:  irm https://claude.ai/install.ps1 | iex
   ```

   Open a new window and type `claude --version` to check it worked.

3. **Go to your folder**: **Claude projects** the first time (to paste the sentence from Step 3), your project folder after that. Mac: type `cd` and a space, drag the folder into the window, and press Enter. Windows: open the folder in File Explorer, type `powershell` in the address bar, and press Enter.
4. Type `claude`. The first time, it opens your browser to sign in.

Recent versions start the terminal in **auto** mode. The kit's settings ask for the asking mode instead; check the mode shown at the bottom, and press `Shift + Tab` to change it. Claude Code has its own dictation in the terminal: type `/voice`, then hold Space and talk. Your audio goes to Anthropic.

Claude Code has built-in commands such as `/resume`, so the kit's skills are named `/start` and `/signoff` to avoid clashes.

Off-computer backup with GitHub

Optional. Make a free **private** repository on github.com and ask Claude to "back this up to GitHub". Sign in through your browser when asked, and never paste passwords or tokens into the chat. Your name and email go with every snapshot (GitHub offers a private "noreply" email for this), and anything once saved stays in the history, even if deleted later. The `private` and `inbox` folders are never included.